Threads, IPC, and Synchronization
CS 111
Summer 2025
Operating System Principles
Peter Reiher

Outline

- Threads
- Interprocess communications
- Synchronization
  - Critical sections
  - Asynchronous event completions

Threads

- Why not just processes?
- What is a thread?
- How does the operating system deal with threads?

Why Not Just Processes?

- Processes are very expensive
  - To create: they own resources
  - To dispatch: they have address spaces
- Different processes are very distinct
  - They cannot share the same address space
  - They cannot (usually) share resources
- Not all programs require strong separation
  - Multiple activities working cooperatively for a single goal
  - Mutually trusting elements of a system

What Is a Thread?

- Strictly a unit of execution/scheduling
  - Each thread has its own stack, PC, registers
  - But other resources are shared with other threads
- Multiple threads can run in a process
  - They all share the same code and data space
  - They all have access to the same resources
  - This makes them cheaper to create and run
- Sharing the CPU between multiple threads
  - User level threads (with voluntary yielding)
  - Scheduled system threads (with preemption)

When Should You Use Processes?

- To run multiple distinct programs
- When creation/destruction are rare events
- When running agents with distinct privileges
- When there are limited interactions and shared resources
- To prevent interference between executing interpreters
- To firewall one from failures of the other

When Should You Use Threads?

- For parallel activities <u>in a single program</u>
- When there is frequent creation and destruction
- When all can run with same privileges
- When they need to share resources
- When they exchange many messages/signals
- When you don’t need to protect them from each other

Thread State and Thread Stacks

- Each thread has its own registers, PS, PC
- Each thread must have its own stack area
- Maximum stack size specified when thread is created
  - A process can contain many threads
  - They cannot all grow towards a single hole
  - Thread creator must know max required stack size
  - Stack space must be reclaimed when thread exits
- Procedure linkage conventions are unchanged

User Level Threads Vs. Kernel Threads

> By now you should be able to deduce the advantages and disadvantages of each

- Kernel threads:
  - An abstraction provided by the kernel
  - Still share one address space
  - But scheduled by the kernel
    - So multiple threads can use multiple cores at once
  - Modern Linux provides kernel threads
- User level threads:
  - Kernel knows nothing about them
  - Provided and managed via user-level library
  - Scheduled by library, not by kernel
- pthread library is a general thread interface

Communications Between Processes

- Even fairly distinct processes may occasionally need to exchange information
- The OS provides mechanisms to facilitate that
  - As it must, since processes can't normally "touch" each other
- These mechanisms are referred to as "inter-process communications"
  - IPC

Goals for IPC Mechanisms

- We look for many things in an IPC mechanism
  - Simplicity
  - Convenience
  - Generality
  - Efficiency
  - Robustness and reliability
- Some of these are contradictory
  - Partially handled by providing multiple different IPC mechanisms

OS Support For IPC

- Provided through system calls
- Typically requiring activity from both communicating processes
  - Usually can't "force" another process to perform IPC
- Usually mediated at each step by the OS
  - To protect both processes
  - And ensure correct behavior

OS IPC Mechanics

- For local processes
- Data is in memory space of sender
- Data needs to get to memory space of receiver
- Two choices:

1. The OS copies the data
2. The OS uses VM techniques to switch ownership of memory to the receiver

IPC: Synchronous and Asynchronous

- Synchronous IPC
  - Writes block until data is sent/delivered/received
  - Reads block until new data is available
  - Very easy for programmers
- Asynchronous operations
  - Writes return when system accepts data
    - No confirmation of transmission/delivery/reception
    - Requires auxiliary mechanism to learn of errors
  - Reads return promptly if no data available
    - Requires auxiliary mechanism to learn of new data
    - Often involves "wait for any of these" operation
  - Much more efficient in some circumstances

Typical IPC Operations

- Create/destroy an IPC channel
- Write/send/put
  - Insert data into the channel
- Read/receive/get
  - Extract data from the channel
- Channel content query
  - How much data is currently in the channel?
- Connection establishment and query
  - Control connection of one channel end to another
  - Provide information like:
    - Who are end-points?
    - What is status of connections?

IPC: Messages vs. Streams

- A fundamental dichotomy in IPC mechanisms
- Streams
  - A continuous stream of bytes
  - Read or write a few or many bytes at a time
  - Write and read buffer sizes are unrelated
  - Stream may contain app-specific record delimiters
- Messages (aka datagrams)
  - A sequence of distinct messages
  - Each message has its own length (subject to limits)
  - Each message is typically read/written as a unit
  - Delivery of a message is typically all-or-nothing
- Each style is suited for particular kinds of interactions

Known by application, not by IPC mechanism

The IPC mechanism knows about these.

IPC and Flow Control

- Flow control: making sure a fast sender doesn't overwhelm a slow receiver
- Queued IPC consumes system resources
  - Buffered in the OS until the receiver asks for it
- Many things can increase required buffer space
  - Fast sender, non-responsive receiver
- Must be a way to limit required buffer space
  - Sender side: block sender or refuse communication
  - Receiving side: stifle sender, flush old data
  - Handled by network protocols or OS mechanism
- Mechanisms for feedback to sender

IPC Reliability and Robustness

- Within a single machine, OS won't accidentally "lose" IPC data
- Across a network, requests and responses can be lost
- Even on single machine, though, a sent message may not be processed
  - The receiver is invalid, dead, or not responding
- And how long must the OS be responsible for IPC data?

Reliability Options

- When do we tell the sender "OK"?
  – When it's queued locally?
  – When it's added to receiver's input queue?
  – When the receiver has read it?
  – When the receiver has explicitly acknowledged it?
- How persistently does the system attempt delivery?
  – Especially across a network
  – Do we try retransmissions? How many?
  – Do we try different routes or alternate servers?
- Do channel/contents survive receiver restarts?
  – Can a new server instance pick up where the old left off?

Some Styles of IPC

- Pipelines
- Sockets
- Shared memory
- There are others we won’t discuss in detail
  - Mailboxes
  - Named pipes
  - Simple messages
  - IPC signals

Pipelines

- Data flows through a series of programs
  - `ls | grep | sort | mail`
  - Macro processor | compiler | assembler
- Data is a simple byte stream
  - Buffered in the operating system
  - No need for intermediate temporary files
- There are no security/privacy/trust issues
  - All under control of a single user
- Error conditions
  - Input: End of File
  - Output: next program failed
- *Simple, but very limiting*

Sockets

- Connections between addresses/ports
  - Connect/listen/accept
  - Lookup: registry, DNS, service discovery protocols
- Many data options
  - Reliable or best effort datagrams
  - Streams, messages, remote procedure calls, …
- Complex flow control and error handling
  - Retransmissions, timeouts, node failures
  - Possibility of reconnection or fail-over
- Trust/security/privacy/integrity
  - We'll discuss these issues later
- *Very general, but more complex*

Shared Memory

- OS arranges for processes to share read/write memory segments (actual shared RAM)
  - Mapped into multiple processes' address spaces
  - Applications must provide their own control of sharing
  - OS is not involved in data transfer
    - They are just memory reads and writes via limited direct execution
    - So <u>very</u> fast
- Simple in some ways
  - Terribly complicated in others
  - The cooperating processes must themselves achieve whatever synchronization/consistency effects they want
- Only works on a local machine

Synchronization

- Making things happen in the “right” order
- Easy if only one set of things is happening
- Easy if simultaneously occurring things don’t affect each other
- Hideously complicated otherwise
- Wouldn’t it be nice if we could avoid it?
- Well, we can’t
  - We must have parallelism

The Benefits of Parallelism

- Improved throughput
  - Blocking of one activity does not stop others
- Improved modularity
  - Separating co… to simpler pieces
- Improved re…
  - The failure … does not stop others
- A better fit to emerging paradigms
  - Client server computing, web based services
  - Our universe is cooperating parallel processes

Kill parallelism and performance goes back to the 1970s

Why Is There a Problem?

- Sequential program execution is easy
  - First instruction one, then instruction two, ...
  - Execution order is obvious and deterministic
- Independent parallel programs are easy
  - If the parallel streams do not interact in any way
- Cooperating parallel programs are hard
  - If the two execution streams are not synchronized
    - Results depend on the order of instruction execution
    - Parallelism makes execution order non-deterministic
    - Results become combinatorially intractable

Synchronization Problems

- Race conditions
- Non-deterministic execution

Race Conditions

- What happens depends on execution order of processes/threads running in parallel
  - Sometimes one way, sometimes another
  - These happen all the time, most don't matter
- But some race conditions affect correctness
  - Conflicting updates (mutual exclusion)
  - Check/act races (sleep/wakeup problem)
  - Multi-object updates (all-or-none transactions)
  - Distributed decisions based on inconsistent views
- Each of these classes can be managed
  - If we recognize the race condition and danger

Non-Deterministic Execution

- Parallel execution makes process behavior less predictable
  - Processes block for I/O or resources
  - Time-slice end preemption
  - Interrupt service routines
  - Unsynchronized execution on another core
  - Queuing delays
  - Time required to perform I/O operations
  - Message transmission/delivery time
- Which can lead to many problems

What Is "Synchronization"?

- True parallelism is too complicated
  - We're not smart enough to understand it
- Pseudo-parallelism may be good enough
  - Mostly ignore it
  - But identify and control key points of interaction
- *Synchronization* refers to that control
- Actually two interdependent problems
  - *Critical section serialization*
  - *Notification of asynchronous completion*
- They are often discussed as a single problem
  - Many mechanisms simultaneously solve both
  - Solution to either requires solution to the other
- They can be understood and solved separately

The Critical Section Problem

- A *critical section* is a resource that is shared by multiple interpreters
  - By multiple concurrent threads, processes or CPUs
  - By interrupted code and interrupt handler
- Use of the resource changes its state
  - Contents, properties, relation to other resources
- Correctness depends on execution order
  - When scheduler runs/preempts which threads
  - Relative timing of asynchronous/independent events

Critical Section Example 1: Updating a File

Process 1 | **Process 2**

```
remove("inventory");                fd = open("inventory",READ);
fd = create("inventory");           count = read(fd,buffer,length);
write(fd,newdata,length);
close(fd);
```

```
    remove("inventory");
    fd = create("inventory");
                                fd = open("inventory",READ);
                                count = read(fd,buffer,length);
    write(fd,newdata,length);
    close(fd);
```

- Process 2 reads an empty file
  - This result could not occur with any sequential execution

Critical Section Example 2: Multithreaded Banking Code

Thread 1

load r1, balance // = 100
load r2, amount1 // = 50
add r1, r2 // = 150
store r1, balance // = 150

load r1, b
load r2, a
add r1, r

Thread 2

load r1, balance // = 100
load r2, amount2 // = 25
sub r1, r2 // = 75
store r1, balance // = 75

The $25 debit was lost!!!

load r1, balance // = 100
load r2, amount2 // = 25
sub r1, r2 // = 75
store r1, balance // = 75

CONTEXT SWITCH!!!

store r1, balance // = 150

amount1 50

balance 150

amount2 25

r1 75

r2 50

Even A Single Instruction Can Contain a Critical Section

thread #1	**thread #2**
counter = counter + 1;	counter = counter + 1;

But what looks like one instruction in C gets compiled to:

```
mov counter, %eax
add $0x1, %eax
mov %eax, counter
```

Three instructions . . .

Why Is This a Critical Section?

thread #1

counter = counter + 1;

thread #2

counter = counter + 1;

This could happen:

```
mov counter, %eax
add $0x1, %eax
```

```
mov counter, %eax
add $0x1, %eax
mov %eax, counter
```

```
mov %eax, counter
```

If counter started at 1, it should end at 3

In this execution, it ends at 2

These Kinds of Interleavings Seem Pretty Unlikely

- To cause problems, things have to happen exactly wrong
- Indeed, that’s true
- But you’re executing a billion instructions per second
- So even very low probability events can happen with frightening frequency
- Often, one problem blows up everything that follows

Critical Sections and Mutual Exclusion

- Critical sections can cause trouble when more than one thread executes them at a time
  - Each thread doing part of the critical section before any of them do all of it
- Preventable if we ensure that only one thread can execute a critical section at a time
- We need to achieve *mutual exclusion* of the critical section
- How?

If one of them is running it, the other definitely isn't!

One Solution: Interrupt Disables

- Temporarily block some or all interrupts
  - No interrupts -> nobody preempts my code in the middle
  - Can be done with a privileged instruction
  - Side-effect of loading new Processor Status Word
- Abilities
  - Prevent Time-Slice End (timer interrupts)
  - Prevent re-entry of device driver code
- Dangers
  - May delay important operations
  - A bug may leave them permanently disabled
  - Won’t solve all sync problems on multi-core machines
    - Since they can have parallelism without interrupts

Downsides of Disabling Interrupts

- Not an option in user mode
  - Requires use of privileged instructions
  - Can be used in OS kernel code, though
- Dangerous if improperly used
  - Could disable preemptive scheduling, disk I/O, etc.
- Delays system response to important interrupts
  - Received data isn’t processed until interrupt serviced
  - Device will sit idle until next operation is initiated
- May prevent safe concurrency

Other Possible Solutions

- Avoid shared data whenever possible
- Eliminate critical sections with atomic instructions
  - Atomic (uninterruptable) read/modify/write operations
  - Can be applied to 1-8 contiguous bytes
  - Simple: increment/decrement, and/or/xor
  - Complex: test-and-set, exchange, compare-and-swap
- Use atomic instructions to implement locks
  - Use the lock operations to protect critical sections
- We'll cover these in more detail in the next class

Conclusion

- Processes are too expensive for some purposes
- Threads provide a cheaper alternative
- Threads can communicate through memory
- Processes need IPC
- Both processes and threads allow parallelism
  - Which is vital for performance
  - But raises correctness issues