

Virtual Machines  
CS 111  
Summer 2025  
Operating System Principles  
Peter Reiher

# Outline

- What is a virtual machine?
- Why do we want them?
- How do virtual machines work?
- Issues in virtual machines

# What Is a Virtual Machine?

- Remember, in CS, “virtual” means “not real”
  - But it looks like it’s real
- So a virtual machine isn’t really a machine
  - But it looks like a machine
- What do we mean by that?
- A *virtual machine* is a software illusion meant to appear to be a real machine
- Virtual machines abbreviated as VMs

# What's That Really Mean?


- We have an actual computer
- We do something in software to make it look like we have multiple computers
  - Or that it's a different kind of computer
  - Making use of the actual computer to do so
- The virtual machine must appear to apps and users to be a real machine

# Graphically, . . .



We implement a virtual server on the real hardware

With a set of virtual components




We have a real server computer

With a real CPU

And real RAM

And real peripherals



# How?

- Use the real hardware to implement the virtual hardware
  - Instructions for the virtual CPU run on the real CPU
  - Real RAM stores the data for virtual RAM
  - A real disk stores data for the virtual disk
  - Etc.
- But to what purpose?

# Why Do We Want Virtual Machines?

- For several reasons
  - Fault isolation
  - Better security
  - To use a different operating system
  - To provide better controlled sharing of the hardware
- Let's consider each reason separately

# Fault Isolation

- Operating systems must never crash
  - Since they take everything down with them
- But crashing a virtual machine's operating system need not take down the entire machine
  - Just the virtual machine
- So our correctness requirements can be relaxed
- Similar advantages for faults that could damage devices
  - They damage the virtual device, not the physical



# Better Security

- The OS is supposed to provide security for processes
- But the OS also provides shared resources
  - Such as the file system and IPC channels
- A virtual machine need not see the real shared resource
- So processes in other virtual machines are harder to reach and possibly damage

# Using a Different Operating System

- Let's say you're running Windows
- But you want to run a Linux executable
- Windows has one system call interface, Linux has a different one
  - So system calls from your Linux executable won't work on Windows
- But if you have a virtual machine running Linux on top of the real machine running Windows . . .
  - Now your application can run fine
  - Assuming you get the virtualization right . . .

# Sharing a Machine's Resources

- In principle, an OS can control how to share resources among processes
- But actually guaranteeing a particular allocation of resources is hard
- It's easier to guarantee an entire virtual machine gets a set allocation of resources
  - So the processes running in it will not steal resources from the other virtual machines
- A very big deal for cloud computing

# How Do We Run Virtual Machines?

- Easiest if the virtual and real machine use the same ISA (Instruction Set Architecture)
  - Tricky and slow, otherwise
  - So the same ISA is the common case
- Basically, rely on limited direct execution
  - Run as much VM activity directly on the CPU as possible
  - When necessary, trap from the VM
- But trap to what?

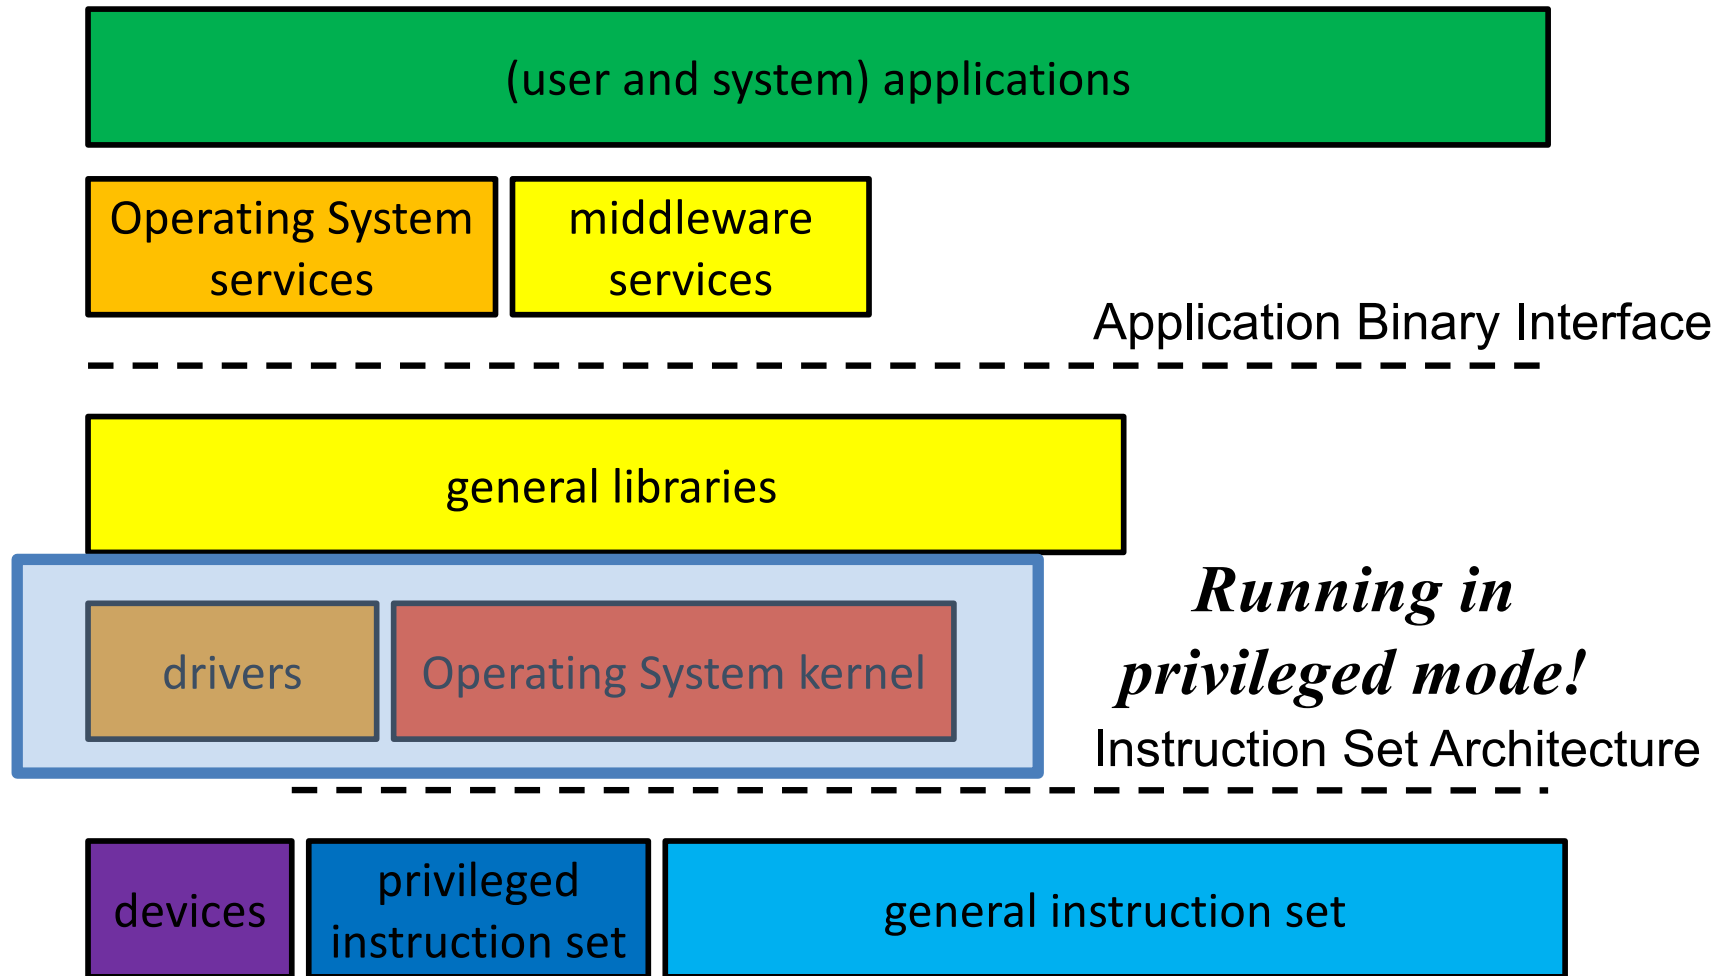
# The Hypervisor

- Also known as the Virtual Machine Monitor (VMM)
- A controller that handles all virtual machines running on a real machine
- When necessary, trap from the virtual machine to the VMM
- It performs whatever magic is necessary
- And then returns to limited direct execution
- Much like a process' system call to an OS

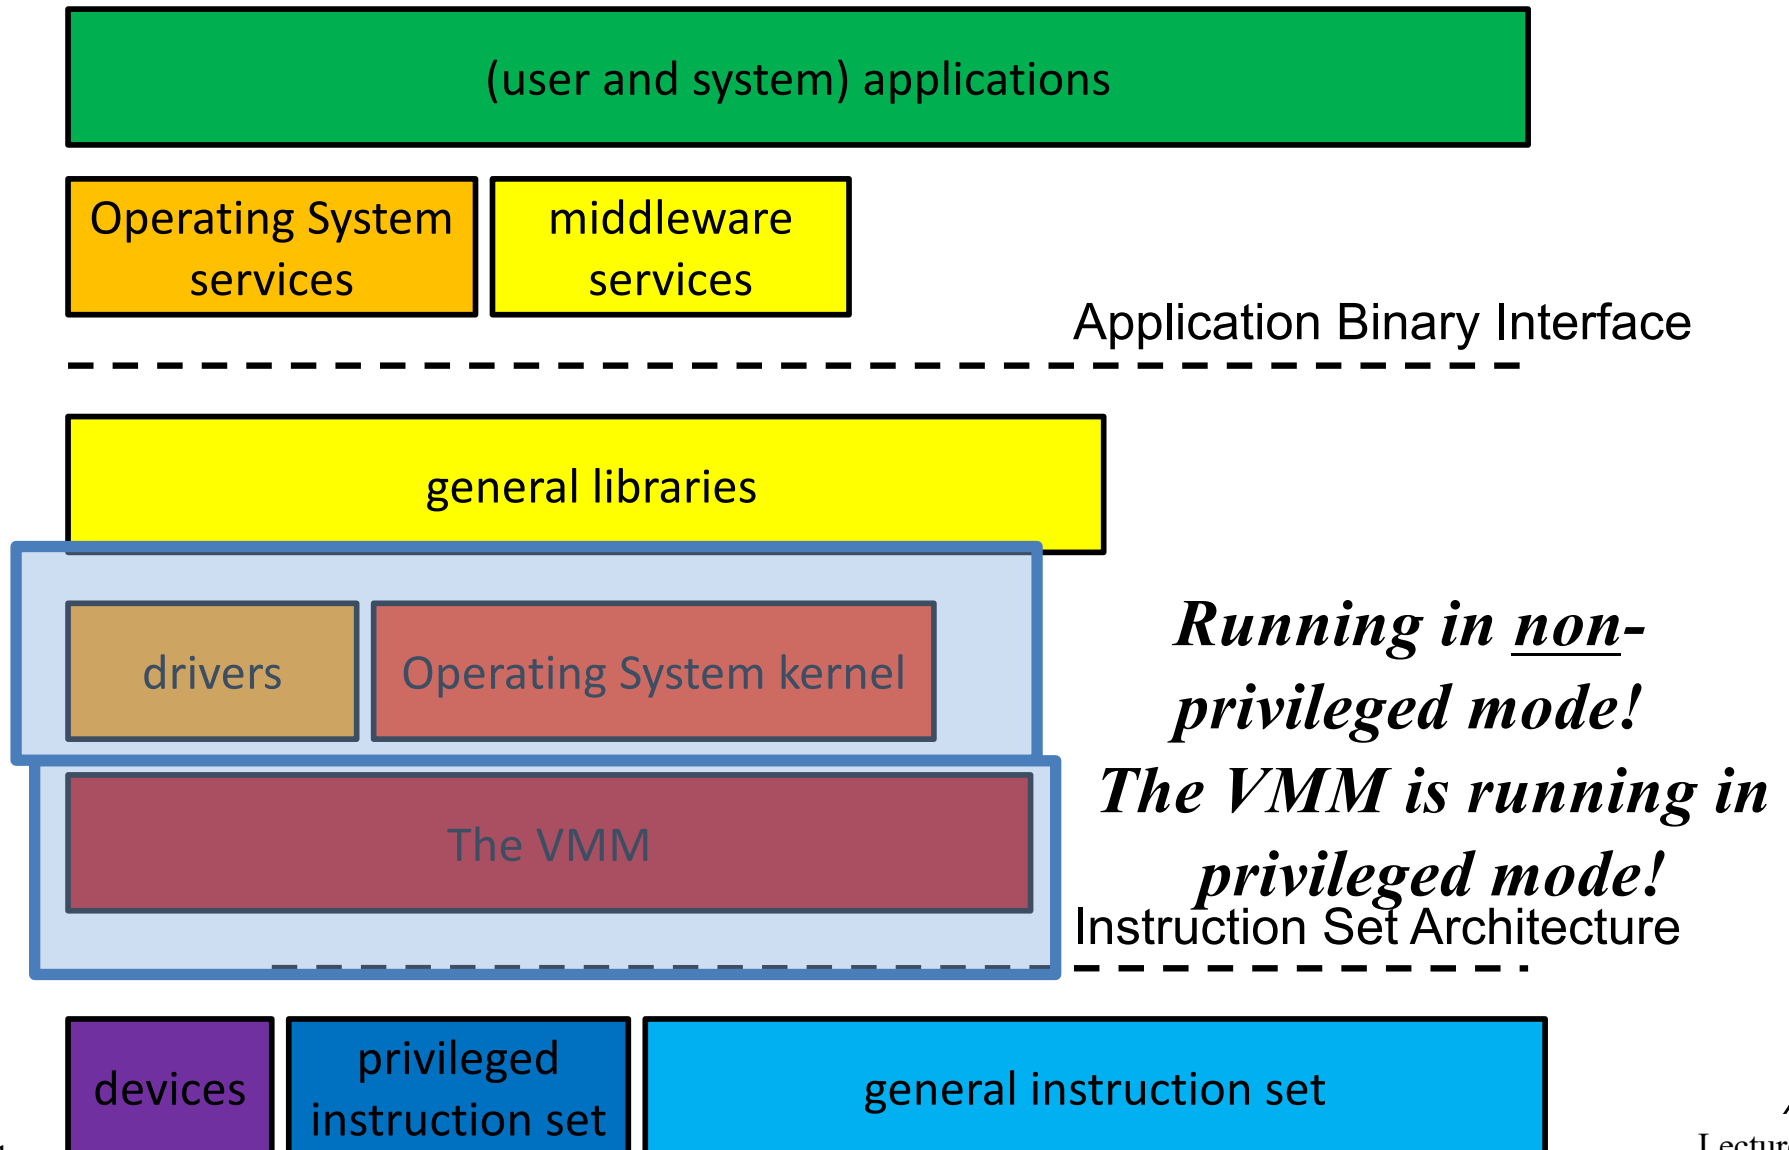
# When Is Trapping to the VMM Necessary?

- Whenever the VM does something privileged
  - Kind of like trapping to the OS when a process wants to do something privileged
- The initial system call instruction will trap to the VMM
- Which will typically forward it to the VM's OS
- But subsequent privileged operations trap back to the VMM

# The Old Architecture



# Architecture With a VMM

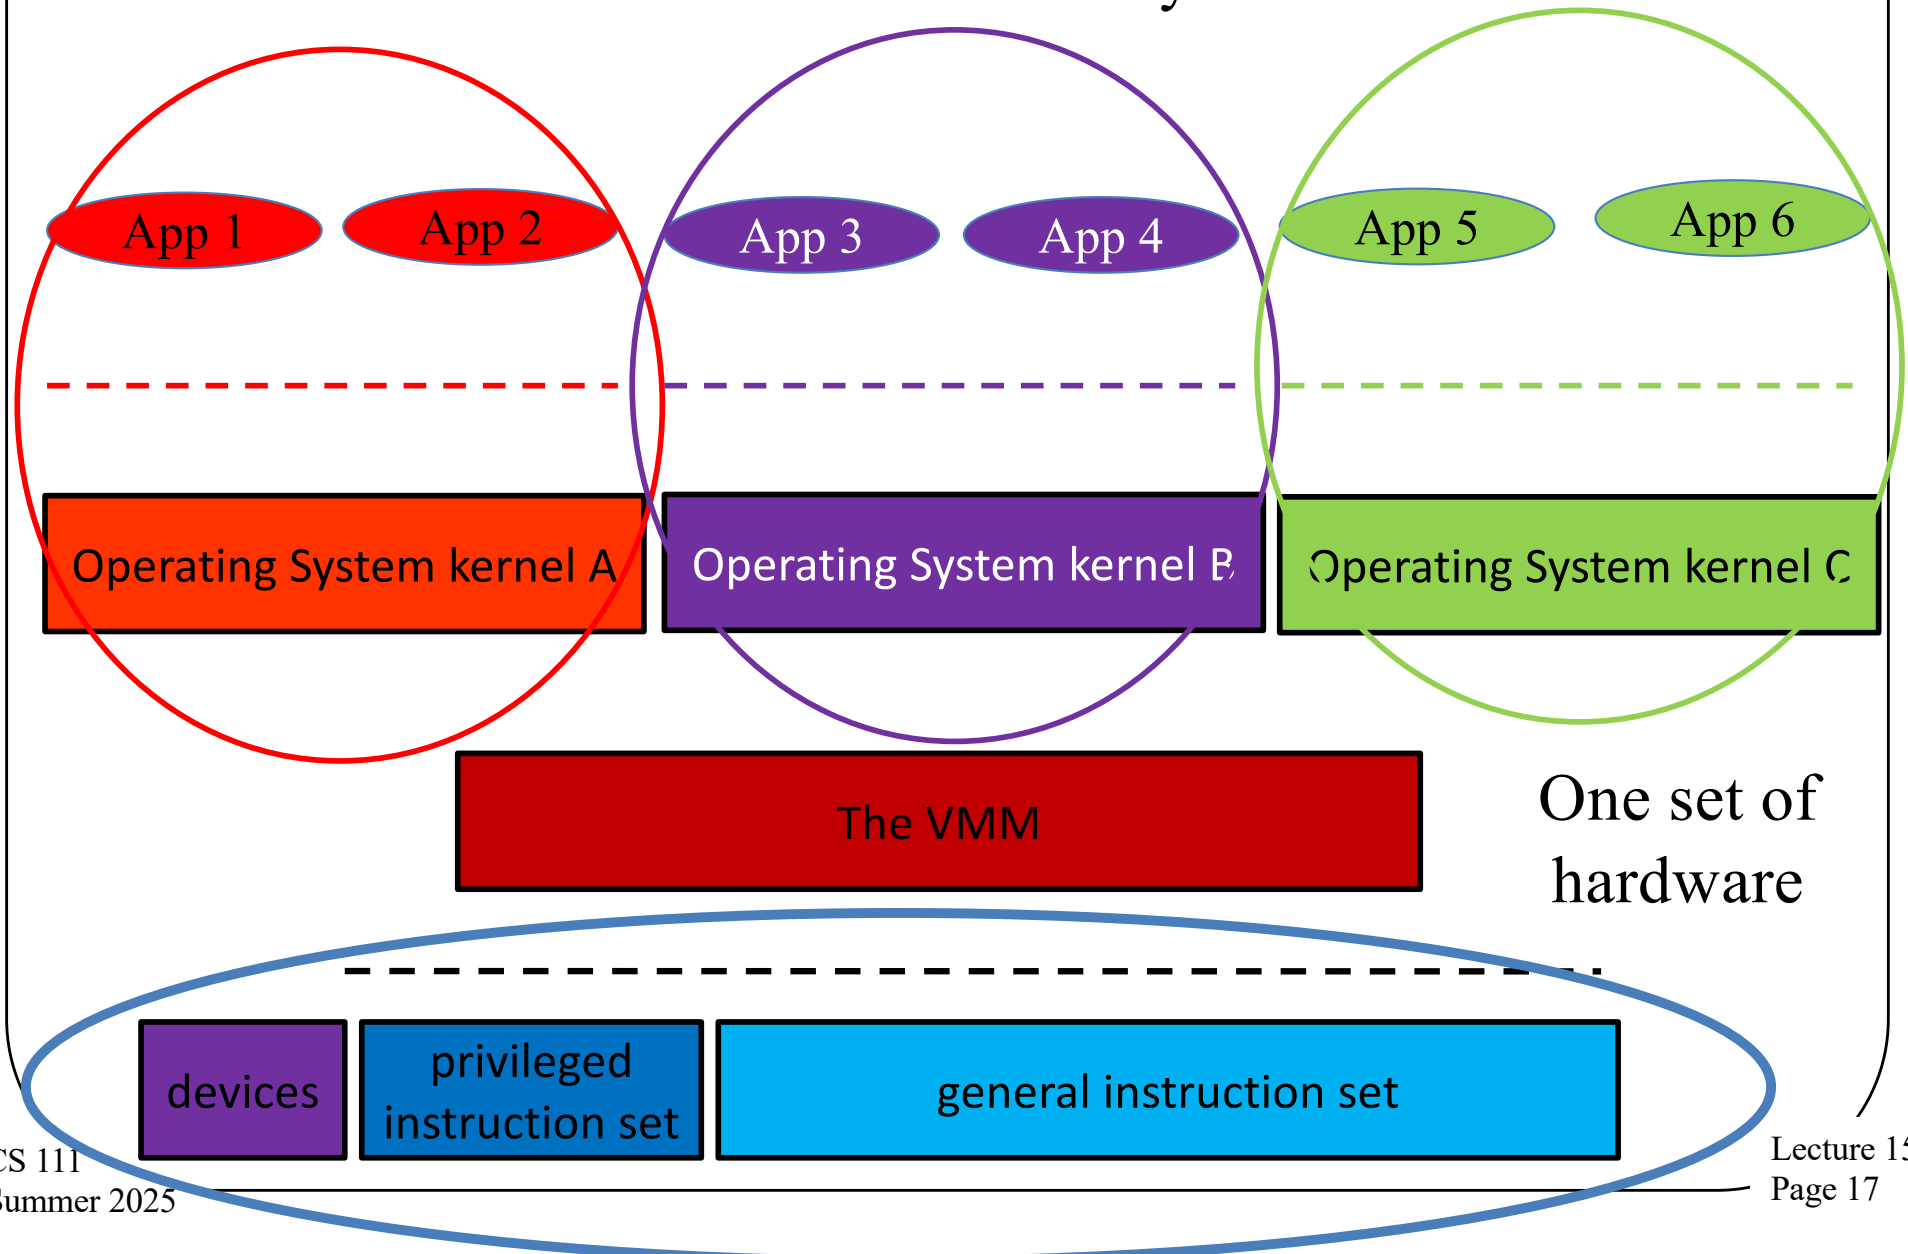




# A More Complex Case

Three virtual machines


With three system call interfaces



# How Do System Calls Work Now?

Using a  
privileged  
instruction

It's sent to  
the VMM  
instead



App 1 makes a  
system call

The virtual  
machine

*Which OS A  
can't perform*



# Yeah, But . . .

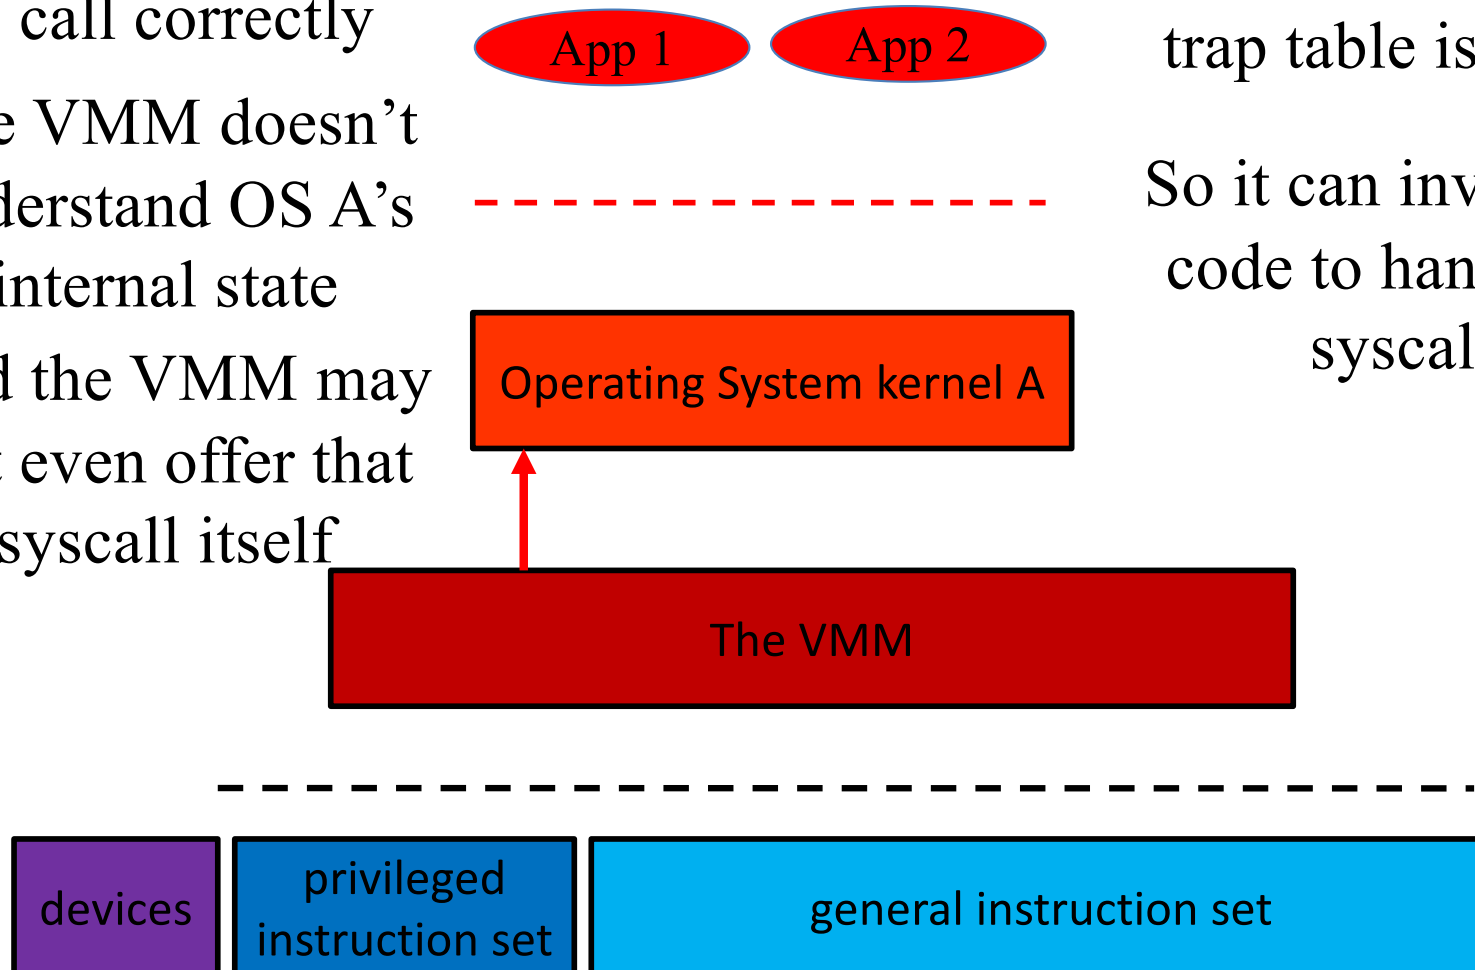
The VMM can't perform the system call correctly

The VMM doesn't understand OS A's internal state

And the VMM may not even offer that syscall itself

But the VMM knows where A's trap table is located

So it can invoke A's code to handle the syscall!




# Yeah, But, Again . . .

If it's a syscall, it  
may need to use  
privileged  
instructions to do its  
work

No problem!

OS A traps when it  
tries to use a  
privileged  
instruction

But OS A can't use  
privileged  
instructions



And the VMM  
catches the trap and  
does the instruction  
for A!




# What's the Point of That?

- If the VMM is going to do the instruction, why not just run A with privilege?
  - So it can do its own instructions
- Well, the VMM might decide not to do the instruction
  - If, for example, it tries to access another VM's memory
- Or the VMM might block VM A and run VM B for a while instead
- The key point: the VMM controls what happens
  - Even though the OS in the VM thinks it is in control

# A Potential Issue

If A is running in non-privileged mode, how can we enforce this interface?



How can we prevent App 1 from messing with A's internal data?

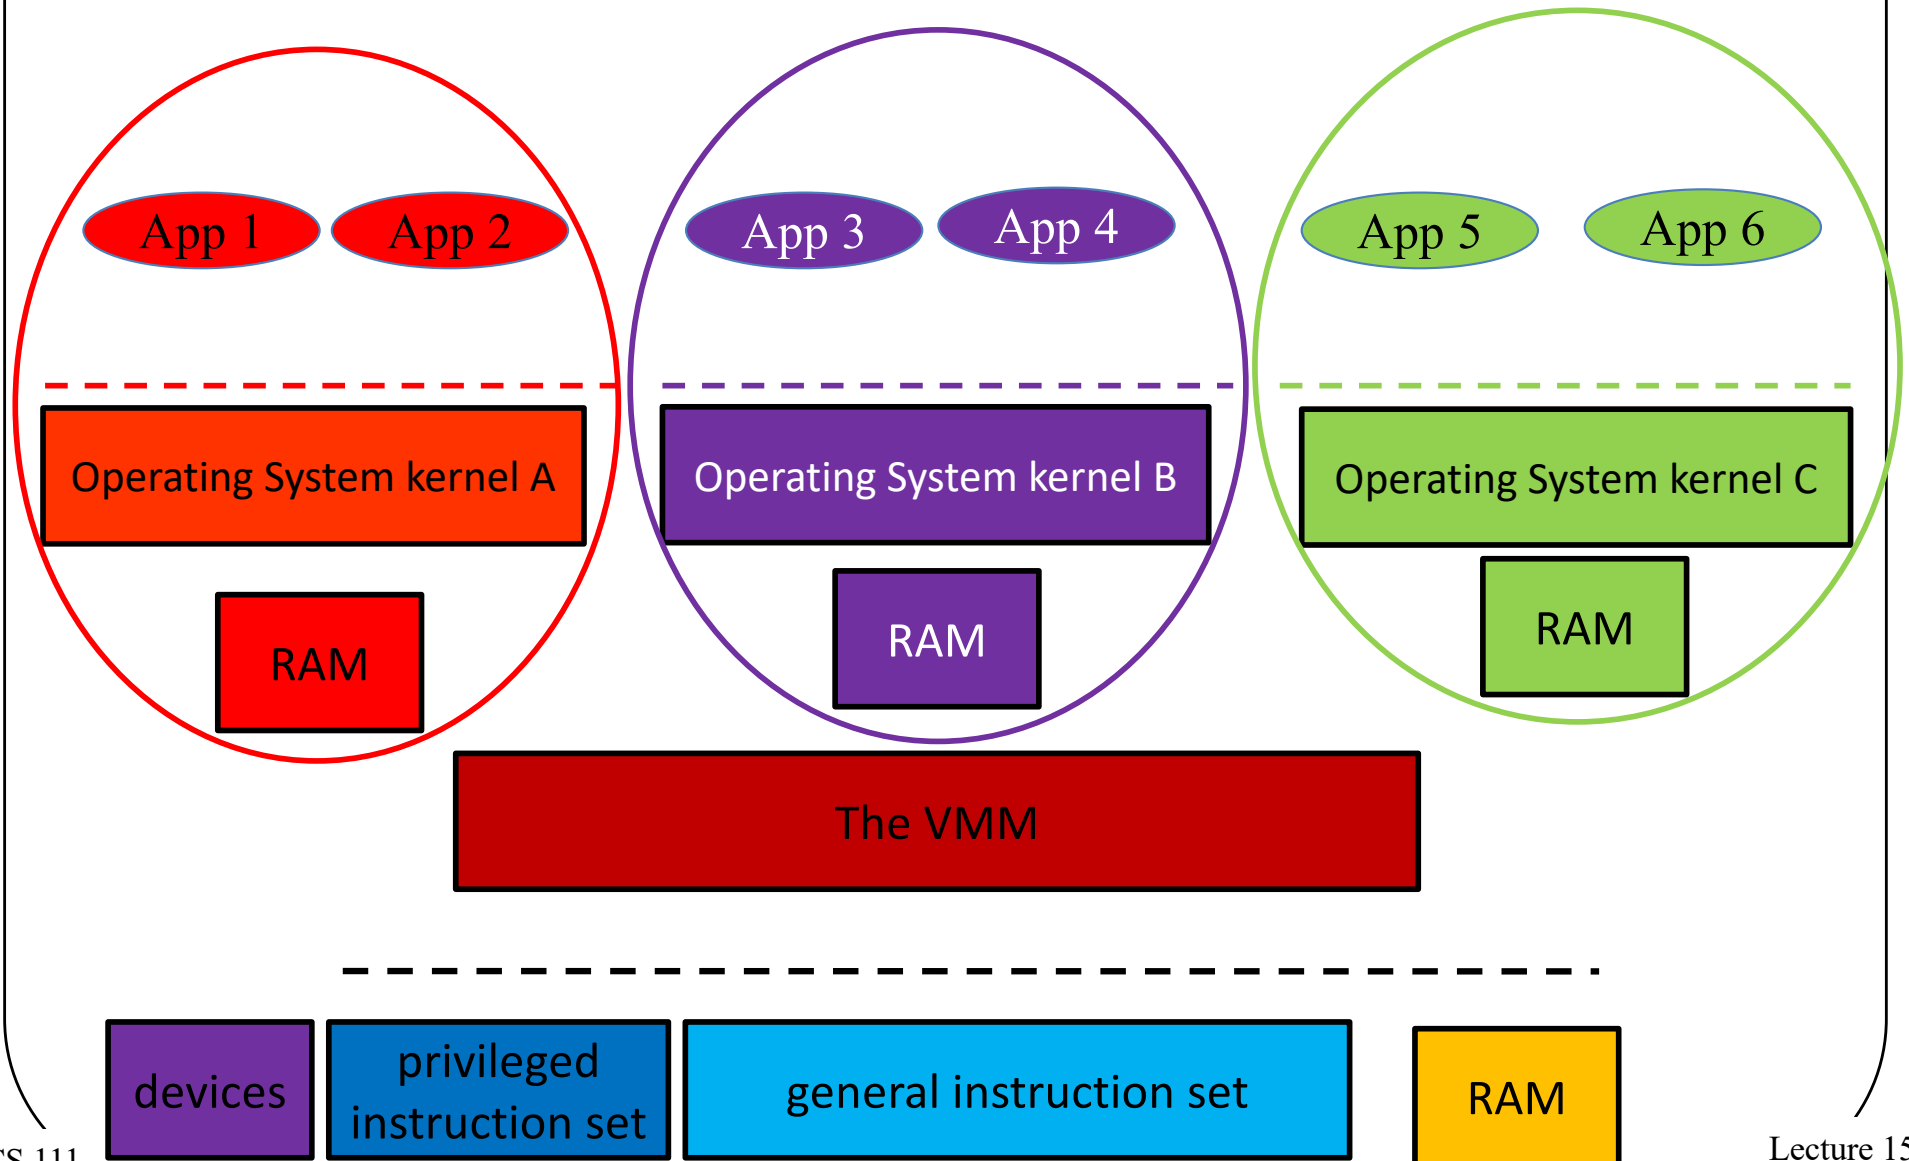
E.g., stop App 1 from killing App 2?



# The Core of the Problem

- OS A thinks it's in control
- OS A believes it's providing segregated virtual memories to App 1 and App 2
- A vital technology for doing so is managing page tables and CPU registers pointing to them
- But OS A has no control over those registers
  - The VMM does
- But the VMM knows nothing of the page tables OS A “controls”

# Virtualizing the Memory








# How To Virtualize Memory

- The virtual OS thinks it has physical memory addresses
  - It provides virtual memory addresses to its processes
  - Handling the virtual-to-physical translations
- The VMM has *machine addresses*
  - Which it translates to physical addresses within a single VM
  - Still using the same paging hardware

# This May Be a Bit Confusing . . .

- Now there are three address types:
  - Virtual addresses  Not RAM addresses!
  - Physical addresses  Also not RAM addresses!
  - Machine addresses  These are RAM addresses!
- In this architecture, “physical addresses” aren’t really physical
  - They aren’t locations on real RAM chips
- Machine addresses are actual locations on real RAM chips

# Three Types of Addresses



A virtual address

*Which translates to . . .*

Which isn't really  
physical



A physical address

*Which translates to . . .*

Which is an actual  
RAM address



A machine address

# For Example

RUNNING  
UNPRIVILEGED

RUNNING  
PRIVILEGED

App 1 issues  
virtual address X

App 1

Causing a TLB miss  
and a trap

The VMM  
invokes OS A

Operating System kernel A

*Since only OS A  
understands App  
1's page table*

The VMM

The VMM  
catches the  
trap

TLB

RAM

# Continuing

And we eventually unwind to run App 1

RUNNING  
UNPRIVILEGED

RUNNING  
PRIVILEGED

App 1

OS A looks up virtual address X in App 1's page table

Operating System kernel A

The VMM translates to the right machine address for X in the TLB and installs it

And tries to install the physical page number for X in the TLB

The VMM

Which causes another trap to the VMM

TLB

RAM

# Looked at Another Way

Some page frame  
actually contains  
page X

App 1

Who knows which  
page frame?

OS A knows that

Operating System kernel A

So the VMM must  
consult OS A to  
perform the  
translation

But the VMM  
doesn't know about  
App 1's address  
space

The VMM

TLB

RAM

The VMM, since it  
controls all page  
frames

# This Can Be (Much) More Complex

- What if the page requested isn't in a page frame?
  - It's somewhere on the flash drive
- Now we not only need to wind and unwind through the VMM
- When the guest OS tries to handle the page fault, that goes through the VMM, too
- And back to the guest OS when the page is delivered to RAM

And what  
about working  
sets?

# Some Implications

- TLB misses are much more expensive
  - Since we'll be moving back and forth from privileged mode to unprivileged
  - Paying overhead costs each time
  - And we'll run more systems code
- We'll need extra paging data structures in the VMM
  - More overhead
- Virtual machines are thus likely to suffer performance penalties



# Making VMs Perform Better

- Adding special hardware
  - Some CPUs have features to make issues of virtualizing the CPU and memory cheaper
- Paravirtualization
  - The basic VM approach assumes the guest OSes in VMs don't know about virtualization
  - If you make some changes to those OSes, they can help make virtualization cheaper

# Virtual Machines and Cloud Computing

- Cloud computing is about sharing hardware among multiple customers
- The cloud provider sells/rents computing services to customers
  - Handling all the difficult issues for them
  - So they can just run their applications
- Cloud providers need lots of customers, to make money
  - Which implies they need lots of hardware

# The Cloud Environment



A warehouse full of vast numbers of machines

Typically tens of thousands

Packed tightly into racks

Connected by high speed internal networks

And connected to the Internet, to allow customers remote access

The expectation is that the environment will run applications for many separate customers at the same time

Many of which might require multiple computers to run properly

With strong guarantees of isolation between customers

# But Why VMs in the Cloud?

- The cloud provider makes the most money by making the most efficient use of the hardware
  - More customers on the same amount of hardware = more profits
- Often, a customer doesn't need the full power of a machine
  - Cloud provider makes more money by using part of that machine for another customer
- But they need strong isolation
- Like that provided by virtual machines . . .

# So . . .

- You run everyone in a virtual machine
- Some customers have many virtual machines to handle their big jobs
  - Perhaps each the only VM on a physical machine
- Some customers' virtual machines share physical machines with other customers' VMs
- Customers' work loads fluctuate
- You want the most efficient packing of VMs onto physical machines possible
  - To maximize profits

# How To Efficiently Place VMs

- There are many physical nodes and many more VMs
- Which do we put where?
- Often reduces to a bin packing algorithm
- Which tends to be NP-hard
  - Where  $n$  may depend on the number of servers and/or VMs considered
  - The more factors considered, the harder to solve
- So estimation techniques are used

# But Things Change . . .

- Customers' jobs vary in their needs
- So a good bin packing for now might be a bad bin packing 10 minutes later
- The cloud controller can readjust as jobs start and end
- But what about long running jobs?
  - Like hosting large web servers?
- Either take away idle VMs
- Or migrate low use VMs to low use hardware

# VMs Aren't Just For Cloud Computing

- As you should know, since your projects use them
- They allow experimentation not easily performed on real hardware
- They allow basic servers to safely divide their resources
- They allow greater flexibility in the software your computer can run



# Conclusion

- Virtual machines are a critical technology for modern computing
- Virtual machines are implemented on real machines
- The key issue is providing each VM the illusion of complete control
- While also providing good performance
- VMs are of special importance in cloud computing